

Linear Regression

Cellphone Dataset

Importing Libraries and Datasets

```
In [ ]: import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn import metrics
import matplotlib.pyplot as plt
```

```
In [ ]: data = pd.read_csv('./Dataset/Cellphone.csv')
data.head()
```

```
Out[ ]:
```

	Product_id	Price	Sale	weight	resoloution	ppi	cpu core	cpu freq	internal mem	ram	RearCam	Front_Cam
0	203	2357	10	135.0	5.2	424	8	1.35	16.0	3.000	13.00	8.0
1	880	1749	10	125.0	4.0	233	2	1.30	4.0	1.000	3.15	0.0
2	40	1916	10	110.0	4.7	312	4	1.20	8.0	1.500	13.00	5.0
3	99	1315	11	118.5	4.0	233	2	1.30	4.0	0.512	3.15	0.0
4	880	1749	11	125.0	4.0	233	2	1.30	4.0	1.000	3.15	0.0

EDA

```
In [ ]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 161 entries, 0 to 160
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Product_id      161 non-null   int64
1   Price           161 non-null   int64
2   Sale            161 non-null   int64
3   weight          161 non-null   float64
4   resoloution     161 non-null   float64
5   ppi             161 non-null   int64
6   cpu core        161 non-null   int64
7   cpu freq        161 non-null   float64
8   internal mem    161 non-null   float64
9   ram             161 non-null   float64
10  RearCam         161 non-null   float64
11  Front_Cam       161 non-null   float64
12  battery         161 non-null   int64
13  thickness       161 non-null   float64
dtypes: float64(8), int64(6)
memory usage: 17.7 KB
```

```
In [ ]: data.isna().sum()
```

```
Out[ ]: Product_id    0
Price              0
Sale              0
weight            0
resolution        0
ppi              0
cpu core          0
cpu freq          0
internal mem      0
ram              0
RearCam           0
Front_Cam         0
battery           0
thickness         0
dtype: int64
```

- 1. Dataset is clean
- 2. We can perform Linear Regression directly

```
In [ ]: data1 = data.iloc[:,[0,2,3,4,5,6,7,8,9,10,11,12,13]]
data1
```

Out[]:

	Product_id	Sale	weight	resolution	ppi	cpu core	cpu freq	internal mem	ram	RearCam	Front_Cam	b
0	203	10	135.0	5.20	424	8	1.350	16.0	3.000	13.00	8.0	
1	880	10	125.0	4.00	233	2	1.300	4.0	1.000	3.15	0.0	
2	40	10	110.0	4.70	312	4	1.200	8.0	1.500	13.00	5.0	
3	99	11	118.5	4.00	233	2	1.300	4.0	0.512	3.15	0.0	
4	880	11	125.0	4.00	233	2	1.300	4.0	1.000	3.15	0.0	
...
156	1206	4638	178.0	5.46	538	4	1.875	128.0	6.000	12.00	16.0	
157	1296	8016	170.0	5.50	534	4	1.975	128.0	6.000	20.00	8.0	
158	856	8809	150.0	5.50	401	8	2.200	64.0	4.000	20.00	20.0	
159	1296	8946	170.0	5.50	534	4	1.975	128.0	6.000	20.00	8.0	
160	1131	9807	202.0	6.00	367	8	1.500	16.0	3.000	21.50	16.0	

161 rows × 13 columns



```
In [ ]: X = data1[['resolution','ppi','cpu core','cpu freq','internal mem','ram','battery']]
```

```
In [ ]: y = data1[:,[1]]
y
```

```
Out [ ]: 

|     | Price |
|-----|-------|
| 0   | 2357  |
| 1   | 1749  |
| 2   | 1916  |
| 3   | 1315  |
| 4   | 1749  |
| ... | ...   |
| 156 | 3551  |
| 157 | 3211  |
| 158 | 3260  |
| 159 | 3211  |
| 160 | 2536  |


```

161 rows × 1 columns

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X,y , random_state=00, test_size=0.30)
```

```
In [ ]: print(X_train.shape)
print(y_train.shape)
print(X_test.shape)
print(y_test.shape)
```

```
(112, 7)
(112, 1)
(49, 7)
(49, 1)
```

Linear Regression Modelling

Modelling on just X , y

```
In [ ]: model = LinearRegression()
model.fit(X,y)
```

```
Out [ ]: ▾ LinearRegression
LinearRegression()
```

```
In [ ]: model.coef_
```

```
Out [ ]: array([[ -2.80094250e+01,  1.08273650e+00,  9.86514030e+01,
                1.93637238e+02,  6.62110199e+00,  1.27977052e+02,
                4.12470925e-02]])
```

```
In [ ]: model.intercept_
```

```
Out [ ]: array([666.93232853])
```

```
In [ ]: model.predict([[5.20,400,5,1.350,16.0,3.000,2310]])
```

```
c:\Users\Prasanna Pandhare\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\b
ase.py:464: UserWarning: X does not have valid feature names, but LinearRegression was fitted
with feature names
  warnings.warn(
```

```
Out[ ]: array([[2294.19477765]])
```

```
In [ ]: model.score(X,y)
```

```
Out[ ]: 0.9331544279402171
```

Performing Modelling on Train-Test data

```
In [ ]: modelFinal = LinearRegression()
```

```
In [ ]: modelFinal.fit(X_train,y_train)
```

```
Out[ ]: ▾ LinearRegression
LinearRegression()
```

```
In [ ]: y_pred = modelFinal.predict(X_test)
```

```
In [ ]: r2_score = metrics.r2_score(y_pred,y_test)
mean_squared_error = metrics.mean_squared_error(y_pred,y_test)
print("R^2 Score: ", r2_score)
print("Mean Squared Error: ", mean_squared_error)
```

```
R^2 Score: 0.9188489800998326
Mean Squared Error: 42911.30679764876
```

```
In [ ]: modelFinal.score(X_test,y_test)
```

```
Out[ ]: 0.9147802517033088
```

Visulization - Graphing

```
In [ ]: plt.scatter(y_test, y_pred, color='blue', label='Actual vs. Predicted')
# plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red', label='Line of
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.title('Actual vs. Predicted Prices for Linear Regression')
plt.legend()
plt.show()
```

Actual vs. Predicted Prices for Linear Regression

